

AI Orchestrator - Complete CLI Usage Guide

This comprehensive guide covers all CLI commands, options, and examples for the AI Orchestrator.



CRITICAL: Virtual Environment Activation

Before running ANY command, you MUST activate the virtual environment:

```
# Navigate to the project directory
cd ~/ai-orchestrator # or your installation directory

# Activate the virtual environment
source venv/bin/activate

# Your prompt should change, e.g.: (venv) user@machine:~$

# Now you can run commands
ai-orchestrator --help
```

Alternative: Use the quick-start script (auto-activates venv):

```
./quick-start.sh --help
./quick-start.sh status
```



Table of Contents

- [Command Overview](#)
- [Testing AI Models](#)
- [Running Tasks](#)
- [Direct Model Queries](#)
- [Status and Configuration](#)
- [Task Analysis](#)
- [API Testing](#)
- [Model Listing](#)
- [Common Workflows](#)
- [Tips and Tricks](#)

Command Overview

Command	Purpose	Example
run	Execute task with orchestration	ai-orchestrator run "Design an API"
ask	Quick query to specific model	ai-orchestrator ask -m anthropic "Hello"
status	Check configuration	ai-orchestrator status
analyze	Preview routing without executing	ai-orchestrator analyze "Build a web app"
test-api	Test API connections	ai-orchestrator test-api
list-models	List available models	ai-orchestrator list-models gemini
init	Initialize .env file	ai-orchestrator init

Testing AI Models

Quick Model Tests

Test each AI model to verify your API keys are working:

```
# Activate venv first!
source venv/bin/activate

# Test OpenAI (ChatGPT)
ai-orchestrator ask -m openai "Say hello in one word"

# Test Anthropic (Claude)
ai-orchestrator ask -m anthropic "Say hello in one word"

# Test Google (Gemini)
ai-orchestrator ask -m gemini "Say hello in one word"

# Test Moonshot (Kimi)
ai-orchestrator ask -m moonshot "Say hello in one word"
```

Test All APIs at Once

```
ai-orchestrator test-api
```

Expected output:

```

Testing OPENAI...
  ✓ SUCCESS - Response: OK
Testing ANTHROPIC...
  ✓ SUCCESS - Response: OK
Testing GEMINI...
  ✓ SUCCESS - Response: OK
Testing MOONSHOT...
  ➔ Skipped (no API key configured)

Summary: 3 passed, 0 failed, 1 skipped

```

Test Specific Model

```

ai-orchestrator test-api -m openai
ai-orchestrator test-api -m anthropic
ai-orchestrator test-api -m gemini
ai-orchestrator test-api -m moonshot

```

Running Tasks

The `run` command is the main entry point for task execution.

Basic Usage

```
ai-orchestrator run "Your task description here"
```

Options

Option	Short	Description
--model	-m	Force specific model (openai/anthropic/gemini/moonshot)
--output	-o	Save output to file
--quiet	-q	Minimal output (response only)
--debug	-d	Show debug information
--env	-e	Path to custom .env file

Examples by Model

OpenAI (ChatGPT) - Architecture & System Design

```
# Auto-routed (detected as architecture task)
ai-orchestrator run "Design a microservices architecture for an e-commerce platform"

# Force OpenAI
ai-orchestrator run -m openai "Create a system design for a real-time chat application"

# Save output to file
ai-orchestrator run -m openai "Design a REST API schema" -o api-design.md
```

Anthropic (Claude) - Coding & Implementation

```
# Auto-routed (detected as coding task)
ai-orchestrator run "Implement a rate limiter in Python with sliding window"

# Force Claude
ai-orchestrator run -m anthropic "Write a React component for a data table with sorting"

# Quiet mode (just the code)
ai-orchestrator run -m anthropic -q "Write a binary search function in Python"
```

Google (Gemini) - Reasoning & Analysis

```
# Auto-routed (detected as reasoning task)
ai-orchestrator run "Analyze the trade-offs between SQL and NoSQL for this use case"

# Force Gemini
ai-orchestrator run -m gemini "Explain why this algorithm has O(n log n) complexity"

# With debug info
ai-orchestrator run -m gemini -d "Compare REST vs GraphQL"
```

Moonshot (Kimi) - Code Review

```
# Auto-routed (detected as review task)
ai-orchestrator run "Review this code for security vulnerabilities"

# Force Kimi
ai-orchestrator run -m moonshot "Audit this function for best practices"
```

Direct Model Queries (ask command)

The `ask` command provides quick, direct access to a specific model.

Syntax

```
ai-orchestrator ask -m <model> "Your prompt"
```

All Models Examples

```
# OpenAI
ai-orchestrator ask -m openai "Explain microservices in 2 sentences"
ai-orchestrator ask -m openai "What's the best way to structure a Python project?"

# Anthropic
ai-orchestrator ask -m anthropic "Write a Python function to reverse a string"
ai-orchestrator ask -m anthropic "How do I fix a memory leak in Node.js?"

# Gemini
ai-orchestrator ask -m gemini "What is 15% of 280?"
ai-orchestrator ask -m gemini "Explain the time complexity of quicksort"

# Moonshot
ai-orchestrator ask -m moonshot "Review this code: def add(a,b): return a+b"
```

Options

Option	Short	Description
--model	-m	Model to use (required)
--quiet	-q	Output only the response
--debug	-d	Show debug information
--env	-e	Path to custom .env file

Debug Mode Example

```
ai-orchestrator ask -m anthropic -d "Test prompt"
```

Output:

```
Debug: Config loaded from default locations
Debug: Prompt: Test prompt
Debug: Creating anthropic client...
Debug: Model name: claude-3-5-sonnet-20241022
Debug: Sending request to anthropic...
Debug: Response received
Debug: Success: True
Debug: Content length: 42
Debug: Tokens used: 15
```

Status and Configuration

Check Current Status

```
ai-orchestrator status
```

Output:

```

AI ORCHESTRATOR v1.0.0

Model Configuration Status

Provider | Model | Status | Specialization |
OpenAI   | gpt-4o-mini | ✓ Configured | Architecture |
Anthropic | claude-3-5-sonnet... | ✓ Configured | Coding |
Gemini   | gemini-2.5-flash | ✓ Configured | Reasoning |
Moonshot | moonshot-v1-8k | ✗ Not configured | Code Review |

✓ 3 model(s) ready for use

```

Initialize New Configuration

```
ai-orchestrator init
```

Creates a new `.env` file in the current directory.

Use Custom Configuration

```

ai-orchestrator status -e /path/to/custom.env
ai-orchestrator run -e /path/to/custom.env "Your task"

```

Task Analysis

Preview how a task would be routed without executing it.

Syntax

```
ai-orchestrator analyze "Your task description"
```

Examples

```

# Simple task
ai-orchestrator analyze "Write a Python function"

# Complex task
ai-orchestrator analyze "Design and implement a REST API with authentication"

```

Output:

Task: Design and implement a REST API with authentication			
Routing Analysis			
#	Task Type	Target	Provider
1	architecture	openai	OpenAI (ChatGPT)
2	coding	anthropic	Anthropic (Claude)

API Testing

Test All Configured Models

```
ai-orchestrator test-api
```

Test Specific Model

```
ai-orchestrator test-api -m openai
ai-orchestrator test-api -m anthropic
ai-orchestrator test-api -m gemini
ai-orchestrator test-api -m moonshot
```

Test All with Custom Config

```
ai-orchestrator test-api -e /path/to/custom.env
```

Model Listing

List available models dynamically from the API.

List Gemini Models

```
ai-orchestrator list-models gemini
```

Output:

Available Gemini Models			
Model Name	Display Name	Input Tokens	Output
gemini-2.5-flash	Gemini 2.5 Flash	1048576	8192
gemini-2.5-pro	Gemini 2.5 Pro	2097152	8192
gemini-1.5-pro	Gemini 1.5 Pro	2097152	8192

✓ Found 3 model(s) available for your API key
 Current configured model: gemini-2.5-flash

To use a different model:
 Add to your .env file: GEMINI_MODEL=<model-name>

Common Workflows

1. Quick Code Generation

```
source venv/bin/activate
ai-orchestrator ask -m anthropic "Write a function to validate email addresses in Python"
```

2. System Design with Output

```
source venv/bin/activate
ai-orchestrator run -m openai "Design a scalable notification system" -o notification-system.md
```

3. Code Review

```
source venv/bin/activate
ai-orchestrator run -m moonshot "Review this code for security: $(cat myfile.py)"
```

4. Multi-Model Task (Auto-Routed)

```
source venv/bin/activate
ai-orchestrator run "Design and implement a user authentication system with tests"
```

5. Batch Testing

```
source venv/bin/activate
# Test all APIs first
ai-orchestrator test-api

# Then run tasks
ai-orchestrator run "Build a todo app backend"
```

6. Debug Failed Requests

```
source venv/bin/activate
ai-orchestrator ask -m anthropic -d "Your prompt"
# Shows detailed debug info for troubleshooting
```

Tips and Tricks

1. Use Quotes for Complex Tasks

```
# Good
ai-orchestrator run "Design a REST API with user authentication"

# Also good (multi-line with quotes)
ai-orchestrator run "Design a REST API with:
- User authentication
- CRUD operations
- Rate limiting"
```

2. Pipe Output to Other Commands

```
# Save to file
ai-orchestrator ask -m anthropic -q "Write a bash script" > script.sh

# Pipe to clipboard (macOS)
ai-orchestrator ask -m anthropic -q "Write a function" | pbcopy
```

3. Use Quiet Mode for Scripts

```
# In a script, use -q for clean output
CODE=$(ai-orchestrator ask -m anthropic -q "Write a function")
echo "$CODE" > myfunction.py
```

4. Check Status Before Running

```
# Quick sanity check
ai-orchestrator status && ai-orchestrator run "Your task"
```

5. Use the Quick-Start Script

```
# Automatically handles venv activation
./quick-start.sh run "Your task"
./quick-start.sh ask -m openai "Hello"
./quick-start.sh status
```

6. Remember Model Specializations

For This...	Use This Model	Command
Architecture	OpenAI	-m openai
Coding	Anthropic	-m anthropic
Analysis	Gemini	-m gemini
Code Review	Moonshot	-m moonshot

Environment Variables

The CLI reads configuration from multiple locations (in order of priority):

1. `~/.config/ai-orchestrator/config.env` (recommended)
2. `~/ai-orchestrator/.env`
3. `./.env` (current directory)

Available Variables

```
# API Keys
OPENAI_API_KEY=sk-....
ANTHROPIC_API_KEY=sk-ant-....
GEMINI_API_KEY=AIza...
MOONSHOT_API_KEY=sk-....

# Model Names (optional overrides)
OPENAI_MODEL=gpt-4o-mini
ANTHROPIC_MODEL=claude-3-5-sonnet-20241022
GEMINI_MODEL=gemini-2.5-flash
MOONSHOT_MODEL=moonshot-v1-8k

# Development Settings
DEBUG=false
LOG_LEVEL=INFO
```

Troubleshooting

Command Not Found

```
# Solution: Activate venv
source venv/bin/activate
```

ModuleNotFoundError

```
# Solution: Activate venv and reinstall
source venv/bin/activate
pip install -e .
```

API Errors

```
# Run with debug flag
ai-orchestrator ask -m anthropic -d "Test"

# Test API connection
ai-orchestrator test-api -m anthropic

# Check status
ai-orchestrator status
```

Empty Response

```
# Try with different prompt
ai-orchestrator ask -m openai "Say hello"

# Check debug output
ai-orchestrator ask -m openai -d "Say hello"
```

See Also

- [README.md](#) (README.md) - Project overview
- [MANUAL_SETUP.md](#) (MANUAL_SETUP.md) - Manual installation guide
- [QUICK_REFERENCE.md](#) (QUICK_REFERENCE.md) - Quick copy-paste commands
- [TROUBLESHOOTING.md](#) (TROUBLESHOOTING.md) - Detailed troubleshooting
- [MODELS.md](#) (MODELS.md) - Available models documentation

Last Updated: February 2026